

Preface

The demand to write this document become a self-evident after several practical instances of deployed data base historians for Ignition based production systems. These systems required a large set of telemetry tags to be stored and managed over a vast time period. The first choice was the MS SQL DB engine, Express or Standard editions. While the Express edition imposes restrictions on CPU usage, db size and lack of built-in scheduler, the Standard edition requires paid license and maintenance plan to be developed and implemented. The alternative options were; the most popular MySQL & PostgreSQL with Timescale extension. Both are open source and available for commercial use with free license. The PostgreSQL with Timescale option was chosen. The document provides a complete set of instructions and steps in order to deploy PostgreSQL DB with Timescale extension as a storage system for Ignition Tag Historian Module. The document does not provide instructions for installing the OS and Ignition software. It is expected the reader knows how to install the OS and Ignition software, and to configure communication driver with automation controller. It is also expected that the reader knows the foundation of networking technologies and protocols.

Introduction

Why PostgreSQL with Timescale? The primary objective to address the “Maintenance” issue. This issue is notorious with MS SQL DB. Log file for MS SQL DB (do not confuse with the data file) will swell over time, and may consume the entire disk space if maintenance plan is not deployed. The secondary objective to address performance degradation. While MS SQL DB engine offers an excellent performance in the beginning, later, as the amount of stored historical data grow, the performance issue begin to emerge. The issue becomes even more prominent if there no maintenance plan deployed. The “later” term in this case is vaguely defined as 1+ year worth of history for a 5000+ tags. How PostgreSQL with Timescale addressing the above issue? The “Maintenance” issue addressed by enabling 2 maintenance jobs; one for data retention and the other for data compression. Note, that these functions are only available for HyperTables, and this is why the Timescale extension for PostgreSQL is imperative. (Timescale extension enables the HyperTable functionality in PostgreSQL DB engine). The “Performance” issue addressed while using overall more efficient method of storing data (saves disk space) and utilization of BRIN(Block Range INdex). The benefit of it (BRIN) becomes prominent in data retrieval when the storage table grows in size (tens or hundreds of GB or TB).

Environment used for the case described depicted in **(Fig.1)**. Both hosts are virtual machines.

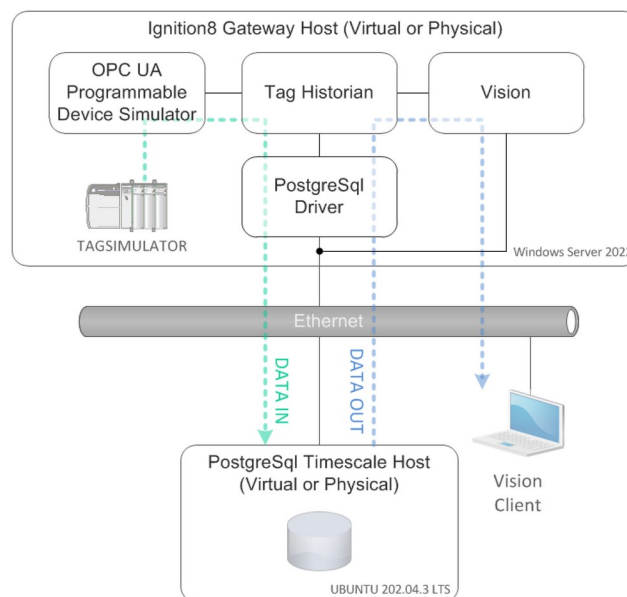


Fig.1

Prerequisites

- Intermediate knowledge of Ignition 8 application.
- General knowledge of operating systems such as Microsoft Windows and UBUNTU.
- General knowledge of network technology.
- Basic-to-Intermediate knowledge of data base technology and SQL.

Credits

- This manual is partly based on <https://docs.timescale.com/self-hosted/latest/install/installation-linux/> by the Timescale. Modifications and additional sections were introduced to address topics specific for integration with Ignition Tag Historian module.
- Virtual environment constructed with free version of ESXi 7.0U3g hypervisor provided by the VMWare.
- PostgreSQL 14.10 license free version provided by the PostgreSQL.
- PgAdmin4 version 8.2 Desktop free license provided by the pgAdmin Development Team.
- Ignition 8.1.37 LTS time limited license provided by the Inductive Automation, LLC.

Instructions (begins with UBUNTU OS host)

It is expected that 2 virtual or physical hosts already deployed, and Ignition software installed on one of them. The host with UBUNTU OS (scheduled for PostgreSQL installation) must have access to the Internet.

- **Checking the UBUNTU OS version** Click the “Settings” > “About” (Fig.2).

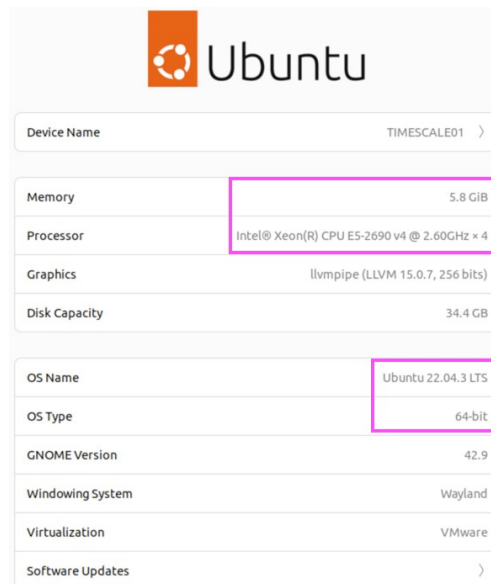


Fig.2

- **Installing the PostgreSQL application in Linux Terminal**

```
sysadmin@TIMESCALE01:~$ sudo apt install gnupg postgresql-common apt-transport-https lsb-release wget
[sudo] password for sysadmin: *****
Reading package lists... Done
Building dependency tree... Done
Reading state information...
...
After this operation, 1,608 kB of additional disk space will be used.
Do you want to continue? [Y/n] Y
Get:1 http://us.archive.ubuntu.com/ubuntu jammy-updates/universe amd64 apt-transport-https all 2.4.11 [1,510 B]
Get:2 http://us.archive.ubuntu.com/ubuntu jammy/main amd64 libcommon-sense-perl amd64 3.75-2build1 [21.1 kB]
...
Removing obsolete dictionary files:
Created symlink /etc/systemd/system/multi-user.target.wants/postgresql.service → /lib/systemd/system/postgresql.service.
Processing triggers for man-db (2.10.2-1) ...
sysadmin@TIMESCALE01:~$ sudo /usr/share/postgresql-common/pgdg/apt.postgresql.org.sh
This script will enable the PostgreSQL APT repository on apt.postgresql.org on
your system. The distribution codename used will be jammy-pgdg.
```

Press **Enter** to continue, or Ctrl-C to abort.

```

Writing /etc/apt/sources.list.d/pgdg.list ...
sysadmin@TIMESCALE01:~$ sudo echo "deb https://packagecloud.io/timescale/timescaledb/ubuntu/ $(lsb_release -c -s)
main" | sudo tee /etc/apt/sources.list.d/timescaledb.list
deb https://packagecloud.io/timescale/timescaledb/ubuntu/ jammy main
sysadmin@TIMESCALE01:~$ wget --quiet -O - https://packagecloud.io/timescale/timescaledb/gpgkey | sudo gpg --dearmor -o
/etc/apt/trusted.gpg.d/timescaledb.gpg
sysadmin@TIMESCALE01:~$ sudo apt update
...
sysadmin@TIMESCALE01:~$ sudo apt install timescaledb-2-postgresql-14
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
...
Processing triggers for libc-bin (2.35-0ubuntu3.6) ...
Processing triggers for man-db (2.10.2-1) ...
sysadmin@TIMESCALE01:~$ sudo timescaledb-tune
Using postgresql.conf at this path:
/etc/postgresql/14/main/postgresql.conf
Is this correct? [(y)es/(n)o]: y
Writing backup to:
/tmp/timescaledb_tune.backup202402012318

shared_preload_libraries needs to be updated
Current:
#shared_preload_libraries = ''
Recommended:
shared_preload_libraries = 'timescaledb'
Is this okay? [(y)es/(n)o]: y
success: shared_preload_libraries will be updated
Tune memory/parallelism/WAL and other settings? [(y)es/(n)o]: y
Recommendations based on 5.78 GB of available memory and 4 CPUs for PostgreSQL 14

Memory settings recommendations
Current:
shared_buffers = 128MB
#effective_cache_size = 4GB
#maintenance_work_mem = 64MB
#work_mem = 4MB
Recommended:
shared_buffers = 1479MB
effective_cache_size = 4437MB
maintenance_work_mem = 75733kB
work_mem = 5048kB
Is this okay? [(y)es/(s)kip/(q)uit]: y
success: memory settings will be updated

Parallelism settings recommendations
Current:
missing: timescaledb.max_background_workers
#max_worker_processes = 8
#max_parallel_workers_per_gather = 2
#max_parallel_workers = 8
Recommended:
timescaledb.max_background_workers = 16
max_worker_processes = 23
max_parallel_workers_per_gather = 2
max_parallel_workers = 4
Is this okay? [(y)es/(s)kip/(q)uit]: y
success: parallelism settings will be updated

WAL settings recommendations
Current:
#wal_buffers = -1
min_wal_size = 80MB
Recommended:
wal_buffers = 16MB
min_wal_size = 512MB
Is this okay? [(y)es/(s)kip/(q)uit]: y
success: WAL settings will be updated

Background writer settings recommendations
Current:
Recommended:
Is this okay? [(y)es/(s)kip/(q)uit]: y
success: background writer settings will be updated

Miscellaneous settings recommendations
Current:
#default_statistics_target = 100
#random_page_cost = 4.0
#checkpoint_completion_target = 0.9
max_connections = 100
#max_locks_per_transaction = 64
#autovacuum_max_workers = 3
#autovacuum_naptime = 1min
#effective_io_concurrency = 1
Recommended:
default_statistics_target = 100
random_page_cost = 1.1
checkpoint_completion_target = 0.9

```

```

max_connections = 75
max_locks_per_transaction = 64
autovacuum_max_workers = 10
autovacuum_naptime = 10
effective_io_concurrency = 256
Is this okay? [(y)es/(s)kip/(q)uit]: y
success: miscellaneous settings will be updated
Saving changes to: /etc/postgresql/14/main/postgresql.conf

```

- Installing psql client with the apt package manager

```

sysadmin@TIMESCALE01:~$ sudo apt-get update
...
sysadmin@TIMESCALE01:~$ sudo apt-get install postgresql-client
Reading package lists... Done
Building dependency tree... Done
...
After this operation, 9,083 kB of additional disk space will be used.
Do you want to continue? [Y/n] Y
Get:1 http://apt.postgresql.org/pub/repos/apt jammy-pgdg/main amd64 postgresql-client-16 amd64 16.1-1.pgdg22.04+1 [1,889 kB]
Get:2 http://apt.postgresql.org/pub/repos/apt jammy-pgdg/main amd64 postgresql-client all 16+256.pgdg22.04+1 [68.9 kB]
Fetched 1,958 kB in 2s (1,296 kB/s)
...
Removing obsolete dictionary files:
sysadmin@TIMESCALE01:~$

```

- Setting up the TimescaleDB extension on UBUNTU system and creating new tsdb database for historian

```

sysadmin@TIMESCALE01:~$ sudo systemctl restart postgresql
sysadmin@TIMESCALE01:~$ systemctl enable postgresql.service
Synchronizing state of postgresql.service with SysV service script with /lib/systemd/systemd-sysv-install.
Executing: /lib/systemd/systemd-sysv-install enable postgresql
sysadmin@TIMESCALE01:~$ sudo -i -u postgres
postgres@TIMESCALE01:~$ psql
psql (16.1 (Ubuntu 16.1-1.pgdg22.04+1), server 14.10 (Ubuntu 14.10-1.pgdg22.04+1))
Type "help" for help.

```

```

postgres=# \password postgres
Enter new password for user "postgres": 
Enter it again: 
postgres=# \q
postgres@TIMESCALE01:~$ psql
psql (16.1 (Ubuntu 16.1-1.pgdg22.04+1), server 14.10 (Ubuntu 14.10-1.pgdg22.04+1))
Type "help" for help.

```

```

postgres=# create user ignition with password '12345';      -- use this password to setup Ignition Database connection
CREATE ROLE
postgres=# create database tsdb;                            -- use this database name to setup Ignition Database connection
CREATE DATABASE
postgres=# grant all privileges on database tsdb to ignition;
GRANT
postgres=# \c tsdb
psql (16.1 (Ubuntu 16.1-1.pgdg22.04+1), server 14.10 (Ubuntu 14.10-1.pgdg22.04+1))
You are now connected to database "tsdb" as user "postgres".
tsdb=# CREATE EXTENSION IF NOT EXISTS timescaledb;
WARNING:
WELCOME TO

```



For more information on TimescaleDB, please visit the following links:

1. Getting started: <https://docs.timescale.com/timescaledb/latest/getting-started>
2. API reference documentation: <https://docs.timescale.com/api/latest>

Note: TimescaleDB collects anonymous reports to better understand and assist our users. For more information and how to disable, please see our docs <https://docs.timescale.com/timescaledb/latest/how-to-guides/configuration/telemetry>.

```

CREATE EXTENSION
tsdb=# \dx

```

Name	Version	Schema	Description
plpgsql	1.0	pg_catalog	PL/pgSQL procedural language
timescaledb	2.13.1	public	Enables scalable inserts and complex queries for time-series data (Community Edition)

(2 rows)

```

tsdb=# \q
postgres@TIMESCALE01:~$ exit
logout
sysadmin@TIMESCALE01:~$

```

- **Modify the PostgreSQL `postgresql.conf` configuration file**

Open the PostgreSQL configuration file “`postgresql.conf`” using text editor. The file located in the “`/etc/postgresql/<version nr>/main`” directory. To open the file in Linux Terminal use the following command:
`sysadmin@TIMESCALE01:~$ sudo nano /etc/postgresql/<version nr>/main/postgresql.conf`
 Find the line `#listen_addresses = 'localhost'` and uncomment it (remove the # character at the beginning of the line). Next, change the value to `listen_addresses = '*'`. This allows PostgreSQL to listen on all available IP addresses. Alternatively, a particular host’s IP address can be specified where the server will listen at.

```
# -----
# PostgreSQL configuration file
# -----
#
# This file consists of lines of the form:
#
#   name = value
#
# (The "=" is optional.)  Whitespace may be used.  Comments are introduced with
# "#" anywhere on a line.  The complete list of parameter names and allowed
# values can be found in the PostgreSQL documentation.
...
#-----
# CONNECTIONS AND AUTHENTICATION
#-----

# - Connection Settings -
#Comment this line:
#listen_addresses = 'localhost'          # what IP address(es) to listen on;
#Add the following:
listen_addresses = '*'

# comma-separated list of addresses;
# defaults to 'localhost'; use '*' for all
# (change requires restart)
# (change requires restart)
# (change requires restart)
# (change requires restart)
port = 5432
max_connections = 75
#superuser_reserved_connections = 3
unix_socket_directories = '/var/run/postgresql' # comma-separated list of directories
```

Save the “`postgresql.conf`” file.

- **Modify the PostgreSQL `pg_hba.conf` configuration file**

Open the PostgreSQL configuration “`pg_hba.conf`” file using text editor. The file located in the “`/etc/postgresql/<version nr>/main`” directory. To open the file from the Linux Terminal use following command:
`sysadmin@TIMESCALE01:~$ sudo nano /etc/postgresql/<version nr>/main/pg_hba.conf`
 Find the following section:

```
# IPv4 local connections:
host    all             all             127.0.0.1/32      md5
And modify it to:
# IPv4 local connections:
host    all             all             0.0.0.0/0         trust

# PostgreSQL Client Authentication Configuration File
# =====
#
# Refer to the "Client Authentication" section in the PostgreSQL
# documentation for a complete description of this file.  A short
# synopsis follows.
#
# This file controls: which hosts are allowed to connect, how clients
# are authenticated, which PostgreSQL user names they can use, which
# databases they can access.  Records take one of these forms:
#
# local      DATABASE  USER  METHOD  [OPTIONS]
# host      DATABASE  USER  ADDRESS METHOD  [OPTIONS]
...
# TYPE  DATABASE      USER                ADDRESS             METHOD
#
# "local" is for Unix domain socket connections only
local   all           all                 peer
# IPv4 local connections:
#Comment this line:
#host    all           all                 127.0.0.1/32       scram-sha-256
#Add the following:
host    all           all                 0.0.0.0/0          trust
...
```

Save the “`pg_hba.conf`” file.

- **Allow tcp port 5432 through the firewall**

To enable tcp traffic on port 5432 through the firewall, from the Linux Terminal use the following command:
`sysadmin@TIMESCALE01:~$ sudo ufw allow 5432/tcp`
 Rules updated
 Rules updated (v6)

- Restart PostgreSQL service

Run the following command from the Linux Terminal to restart PostgreSQL:

```
sysadmin@TIMESCALE01:~$ sudo service postgresql restart
```

- Installing Pgadmin4 utility

```
sysadmin@TIMESCALE01:~$ sudo snap install curl
curl 8.1.2 from Wouter van Bommel (woutervb) installed
sysadmin@TIMESCALE01:~$ curl -fsS https://www.pgadmin.org/static/packages_pgadmin_org.pub | sudo gpg --dearmor -o
/usr/share/keyrings/packages-pgadmin-org.gpg
sysadmin@TIMESCALE01:~$ sudo sh -c 'echo "deb [signed-by=/usr/share/keyrings/packages-pgadmin-org.gpg]
https://ftp.postgresql.org/pub/pgadmin/pgadmin4/apt/$(lsb_release -cs) pgadmin4 main" >
/etc/apt/sources.list.d/pgadmin4.list && apt update'
Hit:1 http://apt.postgresql.org/pub/repos/apt jammy-pgdg InRelease
...
Reading state information... Done
4 packages can be upgraded. Run 'apt list --upgradable' to see them.
N: Skipping acquire of configured file 'main/binary-i386/Packages' as repository 'http://apt.postgresql.org/pub/repos/apt jammy-pgdg
InRelease' doesn't support architecture 'i386'
sysadmin@TIMESCALE01:~$ sudo apt install pgadmin4 -y
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
...
Processing triggers for ufw (0.36.1-4ubuntu0.1) ...
Processing triggers for man-db (2.10.2-1) ...
sysadmin@TIMESCALE01:~$ cd /usr/pgadmin4/bin
sysadmin@TIMESCALE01:/usr/pgadmin4/bin$ sudo apt install pgadmin4
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
pgadmin4 is already the newest version (8.2).
The following packages were automatically installed and are no longer required:
  libflashrom1 libftdi1-2 liblvm13
Use 'sudo apt autoremove' to remove them.
0 upgraded, 0 newly installed, 0 to remove and 4 not upgraded.
sysadmin@TIMESCALE01:/usr/pgadmin4/bin$ cd \
sysadmin@TIMESCALE01:~$ exit
```

- Run pgAdmin4 GUI for the first time

On UBUNTU deck click the “Show Application”  icon. Locate and click the “pgAdmin4” icon .

In pgAdmin4 GUI right-click the “Servers” root and click the Register>Server item. Enter the “Name” then click the “Connection” tab and enter connection parameters. Click the “Save” button (**Fig.3**).

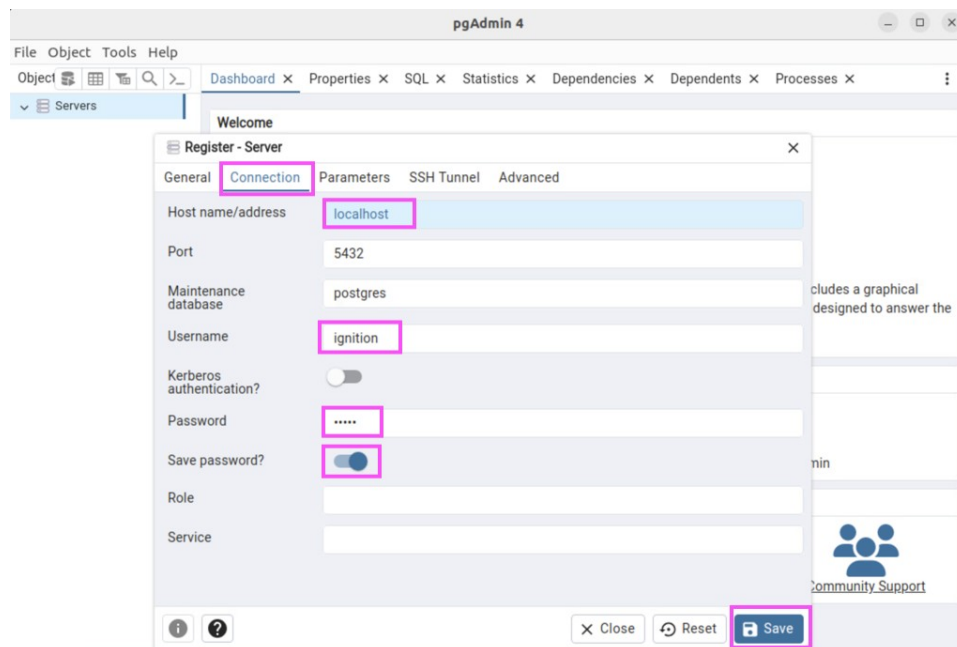


Fig.3

In “Object Explorer” panel expand the tree to reach the “tsdb” node. Click the “tsdb” node (Fig.4).

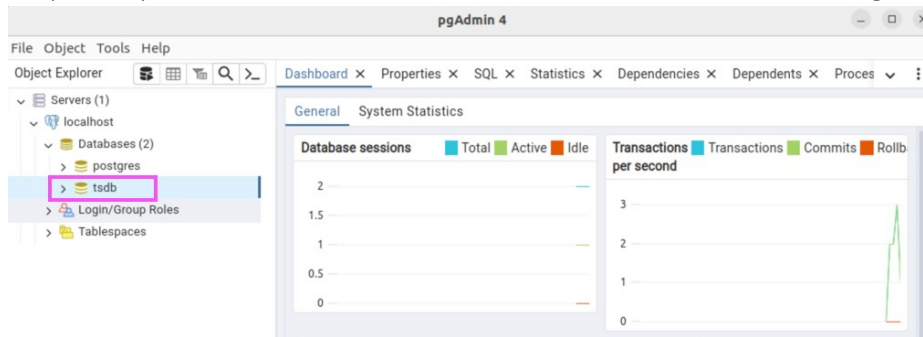


Fig.4

Important! At this point PostgreSQL and pgAdmin4 applications are installed and ready to accommodate Storage System infrastructure for Ignition Tag Historian.

On Ignition 8 Gateway Host:

- Configuring the PostgreSQL driver on Ignition Gateway for the first time

Download (latest) JDBC postgresql driver from <https://jdbc.postgresql.org/> to Ignition gateway host. Copy driver (for instance: postgresql-42.x.y.jar) to the “./Downloads” folder and update driver file in:
Ignition Gateway > Config > DATABASES > Drivers > PostgreSQL > Edit

Note! By default Ignition is using PostgreSQL driver version 9.X that does NOT support Hyper-table features.

Click the “Choose File” button, select driver file from the “./Downloads” folder and click the “Save Changes” button (Fig.5).




Fig.5

Verify the driver file copied to the “C:\Program Files\Inductive Automation\Ignition\user-lib\jdbc” folder. Once driver is updated, configure a new database connection to Postgres tsdb DB using previously defined credentials.
Ignition Gateway > Config > DATABASES > Connections > Create new Database Connection > PostgreSQL > Next

Name = TIMESCALEDB
JDBC Driver = PostgreSQL's
Connect URL = jdbc:postgresql://<ip_address>:5432/tsdb
Username = ignition
Password = 12345 (See: the “Setting up the TimescaleDB extension on UBUNTU system” section)
The rest of the settings are left as default.

Click the “Create New Database Connection” button. Verify created connection status is valid (Fig.6)




Fig.6

- Configure tags history provider

Ignition Gateway > Config > TAGS > History > TIMESCALEDB > edit

Important! The “Data Partitioning” & the “Data Pruning” settings must be disabled! (Fig.7, Fig.8). Both theses functions are to be carried out by the Timescale Hyper-Table extension features. Click the “Save Changes” button.




Fig.7

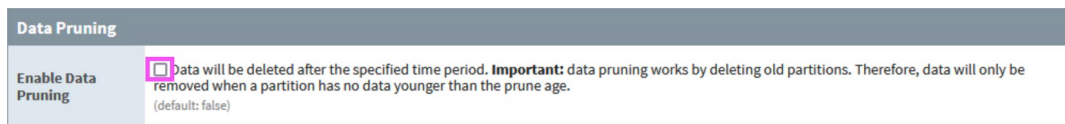


Fig.8

- Create Tags using simulation program

Ignition Gateway > Config > Opcua > Devices > Create new Device > Programmable Device Simulator > Next

Name = TAGSIMULATOR
 Enabled = True
 Repeat Program = True
 Base Rate (ms) = 1000 - period

Click the “Create New Device” button.

Create an empty text file with the name “tags_sim_program.csv” and save it to the “./Download” folder.
 Populate the file with the following data below:

```
Time Interval, Browse Path, Value Source, Data Type
"0", "Ramp/Ramp10", "ramp(0.0, 100.0, 10, true)", "Float"
"0", "Ramp/Ramp20", "ramp(0.0, 100.0, 20, true)", "Float"
"0", "Ramp/Ramp30", "ramp(0.0, 100.0, 30, true)", "Float"
"0", "Ramp/Ramp40", "ramp(0.0, 100.0, 40, true)", "Float"
"0", "Ramp/Ramp50", "ramp(0.0, 100.0, 50, true)", "Float"
"0", "Ramp/Ramp60", "ramp(0.0, 100.0, 60, true)", "Float"
"0", "Sine/Sine10", "sine(0.0, 100.0, 10, true)", "Float"
"0", "Sine/Sine20", "sine(0.0, 100.0, 20, true)", "Float"
"0", "Sine/Sine30", "sine(0.0, 100.0, 30, true)", "Float"
"0", "Sine/Sine40", "sine(0.0, 100.0, 40, true)", "Float"
"0", "Sine/Sine50", "sine(0.0, 100.0, 50, true)", "Float"
"0", "Sine/Sine60", "sine(0.0, 100.0, 60, true)", "Float"
```

In Ignition Gateway > Config > Opcua > Devices > TAGSIMULATOR click the “More” then “edit program”

Set Load Program = Load from CSV, click the “Browse...” button.
 Select the “tags_sim_program.csv” from “./Download” folder and click the “Open”.
 Click the “Load Simulator Program”.
 Verify the list of simulated tags and parameters displayed.
 Click the “Save Program” button. And verify the “TAGSIMULATOR” has status “Running” (Fig.9).




Fig.9

- Create new Ignition Project

Open Ignition designer application and create a new “TIMESACLE” project.

Add device tags to the project: In “Tag Browser” panel click the “+” drop-down list, select the “Browse Devices...” .
 Expand the “Ignition OPC UA Server” to reach the “Ramp” folder. With the “Shift” key select all tags in Ramp folder & add them with the “~” key to default provider window. Perform the same action for all tags in “Sine” folder.
 Click the “OK” button. The “Tag Browser” panel must look like this (Fig.10):




Fig.10

- Create new Tag Group for Timescale

Create "Timescale" tag group. In Tag Browser" panel click the  button, then click the "Edit Group Tags".

Add tag group.
 Name = Timescale
 Mode = Direct
 Rate(ms) = 1000
 Data Mode = Subscribed
 ...
 Min Time Between Samples = 1 sec
 Max Time Between Samples = 1 min

Click the "OK" to save the group.

- Initializing Storage System table infrastructure for Ignition Tag Historian

In Tag Browser" panel double-click the "Sine10" tag. In "Tag Editor" pop-up window select the "History" property from the "Categories" list. Change the following settings as shown:

History Enabled = true
 Storage Provider = TIMESCALEDDB
 Deadband Mode = Percent
 Historical Deadband = 0.01
 Sample Mode = Tag Group
 ...
 Historical Tag Group = Timescale

VERY IMPORTANT! Click the "Apply" button, and within 2-5 sec. set the History Enabled for the "Ramp10" tag back to false. Click the "Apply" button again. At this point there must be absolutely no tags with enabled history provider pointing to the "TIMESCALEDDB".

Disable the history provider "TIMESCALEDDB" in Ignition Gateway > Config > Tags > History >TIMESCALEDDB (**Fig.11**) to reassure that.




Fig.11

On UBUNTU OS PostgreSQL Host:

- Switch back to the PostgreSQL Timescale Historian DB Host

Open the "pgAdmin4" GUI, in "Object Explorer" navigate to:

Servers > Localhost > Databases > tsdb > Schemas > public > Tables

There the Storage System for Ignition Tag Historian has crated a set of tables (**Fig.12**)




Fig.12

On first time historization is enabled for a tag, Ignition creates a table to store historical data with the name "sqlth_1_data". Subsequent setup from scratch on the same Ignition gateway will create another table named "sqlth_2_data", Each subsequent setup will increment the table index by 1. Table index value is managed in Ignition gateway. If the name format other that "sqlth_X_data" is observed, it indicates an error in historical provider configuration. Typically the result of the "Data Partitioning" option left enabled.

Right-click the "sqlth_1_data" table name and click the "Query Tool" item. In "Query" window execute the following query: `select * from public.sqlth_1_data;` Data Output must show several recorded samples of "Sine10" tag (**Fig.13**)

Data Output							
	tagid [PK] integer	intvalue bigint	floatvalue double precision	stringvalue character varying (255)	datevalue timestamp without time zone	dataintegrity integer	t_stamp [PK] bigint
1	1	[null]	5.753056526184082	[null]	[null]	192	1707524912542
2	1	[null]	0.516679584980011	[null]	[null]	192	1707524911541
3	1	[null]	14.225116729736328	[null]	[null]	192	1707524910540
4	1	[null]	41.62026596069336	[null]	[null]	192	1707524909538

Fig.13

- Upgrading the “sqlth_1_data” table to a Hyper-Table

Open Linux Terminal and execute the following commands:

```
sysadmin@TIMESCALE01:~$ sudo -i -u postgres
```

```
postgres@TIMESCALE01:~$ psql
```

```
psql (16.1 (Ubuntu 16.1-1.pgdg22.04+1), server 14.10 (Ubuntu 14.10-1.pgdg22.04+1))
Type "help" for help.
```

```
postgres=# \c tsdb
```

```
psql (16.1 (Ubuntu 16.1-1.pgdg22.04+1), server 14.10 (Ubuntu 14.10-1.pgdg22.04+1))
You are now connected to database "tsdb" as user "postgres".
```

```
tsdb=# select * from sqlth_1_data;
```

tagid	intvalue	floatvalue	stringvalue	datevalue	dataintegrity	t_stamp
1		5.753056526184082			192	1707524912542
1		0.516679584980011			192	1707524911541
1		14.225116729736328			192	1707524910540
1		41.62026596069336			192	1707524909538

(4 rows)

```
tsdb=# SELECT * FROM create_hypertable('sqlth_1_data','t_stamp', if_not_exists => True, chunk_time_interval => 86400000, migrate_data => True);
```

WARNING: column type "character varying" used for "stringvalue" does not follow best practices

HINT: Use datatype TEXT instead.

WARNING: column type "timestamp without time zone" used for "datevalue" does not follow best practices

HINT: Use datatype TIMESTAMPTZ instead.

NOTICE: migrating data to chunks

DETAIL: Migration might take a while depending on the amount of data.

hypertable_id	schema_name	table_name	created
1	public	sqlth_1_data	t

(1 row)

Note! Chunks time interval is set to 86400000 ms OR 24 hrs.

```
tsdb=# ALTER TABLE sqlth_1_data SET (timescaledb.compress, timescaledb.compress_orderby = 't_stamp DESC', timescaledb.compress_segmentby = 'tagid');
```

ALTER TABLE

```
tsdb=# CREATE OR REPLACE FUNCTION unix_now() returns BIGINT LANGUAGE SQL STABLE as $$ SELECT (extract(epoch from now())*1000)::bigint $$;
```

CREATE FUNCTION

```
tsdb=# SELECT set_integer_now_func('sqlth_1_data', 'unix_now');
```

```
set_integer_now_func
```

(1 row)

```
*****
*
* set_integer_now_func()
*
* This function is only relevant for hypertables with integer (as opposed to TIMESTAMP/TIMESTAMPTZ/DATE)time values.
* For such hypertables, it sets a function that returns the now() value (current time) in the units of the time column.
* This is necessary for running some policies on integer-based tables. In particular, many policies only apply to chunks
* of a certain age and a function that returns the current time is necessary to determine the age of a chunk.
*****
```

```
tsdb=# SELECT add_compression_policy('sqlth_1_data', BIGINT '2678400000');
```

```
add_compression_policy
```

1000	- Note! Returns JOB ID added to the

(1 row)

```
*****
*
* add_compression_policy()
*
* Allows to set a policy by which the system compresses a chunk automatically in the background after it reaches a given
* age. Compression policies can only be created on hypertables or continuous aggregates that already have compression
* enabled. To set timescaledb.compress and other configuration parameters for hypertables, use the ALTER TABLE command.
* To enable compression on continuous aggregates, use the ALTER MATERIALIZED VIEW command. To view the policies that you
* set or the policies that already exist, see informational views. Important! For hypertables with integer-based
* timestamps: the time interval should be an integer type (this requires the set_integer_now_func() to be set).
*****
```

The above example of add_compression_policy command compresses a chunk automatically in the background after it reaches age of 2678400000 ms or 31 day. (Typical for Ignition Historian with 1 year data retention policy).

```
tsdb=# SELECT add_retention_policy('sqlth_1_data', BIGINT '31622400000');
```

```
add_retention_policy
```

1001

(1 row)

```
*****
*
* add_retention_policy()
*
* Create a policy to drop chunks older than a given interval of a particular hypertable or continuous aggregate on a
* schedule in the background. For more information, see the drop_chunks section. This implements a data retention policy
* and removes data on a schedule. Only one retention policy may exist per hypertable.
*****
```

The above example of add_retention_policy() command (Typical for Ignition Historian with 1 year (366 days) data retention policy).

Set the “fixed_schedule” flag to TRUE and set specific “next_start” schedules for both jobs.
 Note! In this case the “retention” job 1001 will execute every day at 20:10 (UTC -6 hrs), and the “compression” job 1000 will follow 5 min later at 20:15 (UTC -6 hrs).

```
tsdb=# SELECT alter_job(1000, fixed_schedule => true, initial_start => '2024-02-08 20:15:00.000000-06', next_start => '2024-02-09 20:15:00.000000-06');
```

```
tsdb=# SELECT alter_job(1001, fixed_schedule => true, initial_start => '2024-02-08 20:10:00.000000-06', next_start => '2024-02-09 20:10:00.000000-06');
```

```
*****
*           fixed_schedule
*
*           Set to FALSE if you want the next start of a job to be determined as its last finish time plus the schedule interval.
*           Set to TRUE if you want the next start of a job to begin schedule_interval after the last start.
*****
```

- Checking hypertable maintenance jobs are configured correctly

```
tsdb=# SELECT job_id, application_name, schedule_interval, retry_period, scheduled, fixed_schedule, config, next_start, hypertable_name FROM timescaledb_information.jobs where hypertable_name = 'sqlth_1_data';
```

job_id	proc_name	schedule_interval	scheduled	fixed_schedule	next_start	hypertable_name
1000	policy_compression	1 day	t	t	2024-02-09 20:15:00-06	sqlth_1_data
1001	policy_retention	1 day	t	t	2024-02-09 20:10:00-06	sqlth_1_data

(2 rows)

- Verifying hypertable maintenance jobs execution statistics after the first scheduled run (can be done later)

```
tsdb=# SELECT job_id, last_successful_finish, last_run_status, job_status, next_start, total_runs, total_failures FROM timescaledb_information.job_stats WHERE hypertable_name = 'sqlth_1_data';
```

job_id	last_successful_finish	last_run_status	job_status	next_start	total_runs	total_failures
1000	2024-02-09 20:15:00.023025-06	Success	Scheduled	2024-02-10 20:15:00-06	1	0
1001	2024-02-09 20:10:00.017209-06	Success	Scheduled	2024-02-10 20:10:00-06	1	0

(2 rows)

- Create BRIN (Block Range Index) for the “sqlth_1_data” hypertable

```
tsdb=# CREATE INDEX sqlth_1_data_tagid_t_stamp_brinx ON sqlth_1_data USING BRIN (tagid, t_stamp) with (pages_per_range=64, autosummarize = on);
```

CREATE INDEX

```
tsdb=# SELECT indexname, indexdef FROM pg_indexes WHERE tablename = 'sqlth_1_data';
```

indexname	indexdef
sqlth_1_data_pkey	CREATE UNIQUE INDEX sqlth_1_data_pkey ON public.sqlth_1_data USING btree (tagid, t_stamp)
sqlth_1_datat_stampndx	CREATE INDEX sqlth_1_datat_stampndx ON public.sqlth_1_data USING btree (t_stamp)
sqlth_1_data_tagid_t_stamp_brinx	CREATE INDEX sqlth_1_data_tagid_t_stamp_brinx ON public.sqlth_1_data USING brin (tagid, t_stamp) WITH (pages_per_range='64')

(3 rows)

Note! Indexes “sqlth_1_data_pkey” & “sqlth_1_datat_stampndx” are created automatically on Storage System table infrastructure for Ignition Tag Historian Initialization.

- Optionally, - disable the “sqlth_1_datat_stampndx” index

```
tsdb=# UPDATE pg_index SET indisvalid = false WHERE indexrelid = 'sqlth_1_datat_stampndx'::regclass;
```

UPDATE 1

- Reindex the “sqlth_1_table” hypertable

```
tsdb=# REINDEX (VERBOSE) TABLE public.sqlth_1_data;
```

```
INFO: index "l1_sqlth_1_data_pkey" was reindexed
DETAIL: CPU: user: 0.00 s, system: 0.00 s, elapsed: 0.00 s
INFO: index "_hyper_l1_chunk_sqlth_1_datat_stampndx" was reindexed
DETAIL: CPU: user: 0.00 s, system: 0.00 s, elapsed: 0.00 s
INFO: index "_hyper_l1_chunk_sqlth_1_data_tagid_t_stamp_brinx" was reindexed
DETAIL: CPU: user: 0.00 s, system: 0.00 s, elapsed: 0.00 s
...
INFO: index "3_3_sqlth_1_data_pkey" was reindexed
DETAIL: CPU: user: 0.07 s, system: 0.03 s, elapsed: 0.11 s
INFO: index "_hyper_l3_chunk_sqlth_1_datat_stampndx" was reindexed
DETAIL: CPU: user: 0.04 s, system: 0.01 s, elapsed: 0.06 s
INFO: index "_hyper_l3_chunk_sqlth_1_data_tagid_t_stamp_brinx" was reindexed
DETAIL: CPU: user: 0.04 s, system: 0.00 s, elapsed: 0.04 s
REINDEX
```

- Optional (requires some historical data accumulated) example to demonstrate BRIN index in action

```
tsdb=# EXPLAIN (ANALYZE) SELECT * FROM public.sqlth_1_data WHERE tagid =1 and t_stamp between 1707524933560 and 1708525044573;
```

```

QUERY PLAN
-----
Bitmap Heap Scan on _hyper_1_3_chunk (cost=9.84..1896.22 rows=25903 width=556) (actual time=0.053..22.420 rows=25653 loops=1)
  Recheck Cond: ((tagid = 1) AND (t_stamp >= '1707524933560'::bigint) AND (t_stamp <= '1708525044573'::bigint))
  Rows Removed by Index Recheck: 150743
  Heap Blocks: lossy=1298
-> Bitmap Index Scan on _hyper_1_3_chunk_sqlth_1_data_tagid_t_stamp_brinx (cost=0.00..3.36 rows=33622 width=0) (actual time=0.040..0.040 rows=12980 loops=1)
   Index Cond: ((tagid = 1) AND (t_stamp >= '1707524933560'::bigint) AND (t_stamp <= '1708525044573'::bigint))
Planning Time: 0.557 ms
Execution Time: 23.208 ms
(8 rows)

```

- Enable autovacuum with custom parameters for the “sqlth_1_data” hypertable end exit psql and terminal

```
tsdb=# ALTER TABLE sqlth_1_data SET (autovacuum_enabled = on, autovacuum_analyze_scale_factor = 0.01, autovacuum_vacuum_scale_factor = 0.05, autovacuum_analyze_threshold = 500, autovacuum_vacuum_threshold = 1000);
ALTER TABLE
```

```
tsdb=# \q
postgres@TIMESCALE01:~$ exit
sysadmin@TIMESCALE01:~$ exit
```

On Ignition 8 Gateway Host:

- Re-enable the “TIMESCALE” historical provider

Re-enable the history provider “TIMESCALEDB” in Ignition Gateway > Config > Tags > History >TIMESCALEDB (Fig.14).


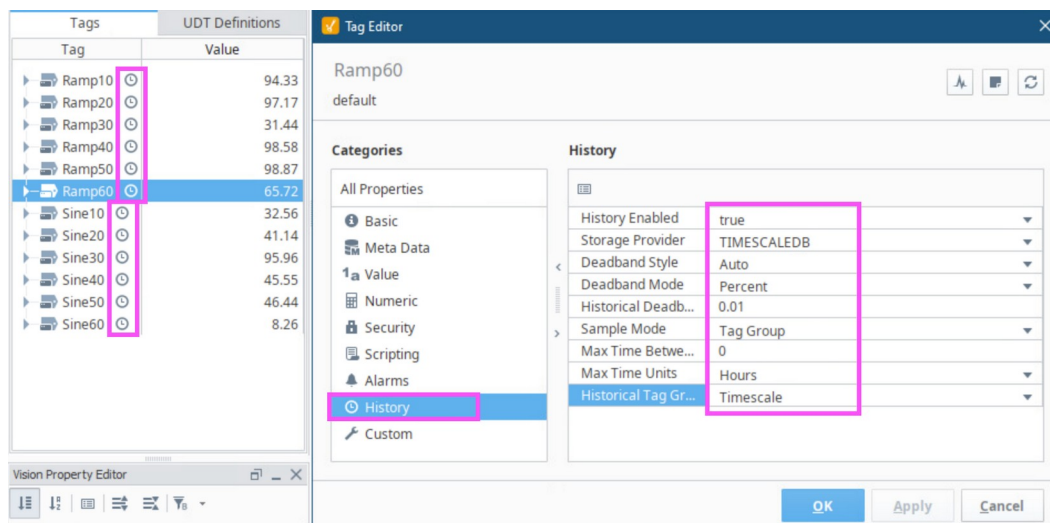
<div>  Successfully updated Historical Tag Provider "TIMESCALEDB" </div>				
Provider Name	Enabled	Type	Description	Status
TIMESCALEDB	true	Datasource History Provider		Running edit

Fig.14

- On Ignition8 host host open the “TIMESACLE” project and finish history setup for simulated all tags.

Enable history for all remaining tags in the “TIMESCALE” project using Storage Provider = TIMESCALEDB and Historical Tag Group = Timescale (Fig.15).



The screenshot shows the Ignition Gateway interface. On the left, a list of tags is displayed, with 'Ramp60' selected. The 'Tag Editor' window is open for 'Ramp60'. In the 'History' tab, the 'History Enabled' checkbox is checked, and the 'Storage Provider' is set to 'TIMESCALEDB'. The 'Historical Tag Group' is set to 'Timescale'. The 'Deadband Style' is set to 'Auto', 'Deadband Mode' is 'Percent', and 'Historical Deadband' is '0.01'. The 'Sample Mode' is 'Tag Group', 'Max Time Between Samples' is '0', 'Max Time Units' is 'Hours', and 'Historical Tag Group' is 'Timescale'. The 'History' tab is highlighted in the left sidebar.

Fig.15

- Create a vision startup page with historical data trends view

Open Ignition designer application and create a startup page with Easy Chart component or similar to allow historical data trending of all 12 tags. In provided example the more versatile trending solution based on Ad Hoc Trends by Matthew Raybourn was used. It can be downloaded from the Ignition Exchange site: <https://inductiveautomation.com/exchange/46/overview>

- Open the TIMESCALE project in Ignition Vision Client Launcher

Example of Ad Hoc Trends component (near real-time) trending TIMESCALE DB historical data for all 12 simulated tags at the same time (**Fig.16**).

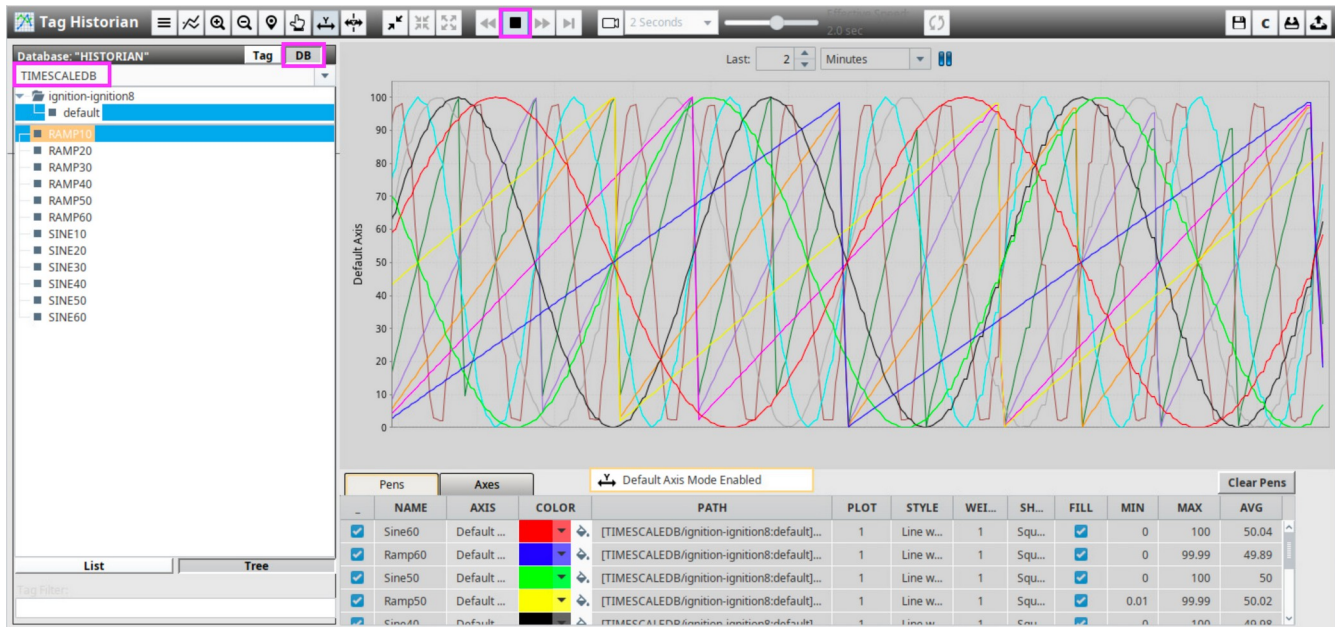


Fig.16

Enjoy!

APPENDIX A

Useful queries for tags sample data, formats, and historical tables infrastructure

Time stamps conversion examples to bigint from dt and from bigint to dt:

```
tsdb=# SELECT now() AS Now_dt, CAST (extract(epoch from now()) *1000 AS BIGINT) AS Now_bigint
, to_timestamp(MAX(t_stamp/1000)) AS MAX_dt, MAX(t_stamp) AS MAX_bigint FROM public.sqlth_1_data;
      now_dt              | now_bigint              | max_dt              | max_bigint
-----+-----+-----+-----
 2024-02-12 17:27:42.7403-06 | 1707780462740          | 2024-02-12 17:27:40-06 | 1707780460382
(1 row)
```

Extract tag's historical data samples with tagname and t_stamp in dt format:

```
tsdb=# SELECT T1.tagid, T2.tagpath, T2.datatype, T1.floatvalue, T1.dataintegrity, T1.t_stamp,
to_timestamp(T1.t_stamp/1000) t_stamp_dt
from public.sqlth_1_data T1, public.sqlth_te T2 where T1.tagid = T2.id AND T2.tagpath ='sine60'
ORDER BY T1.t_stamp DESC FETCH FIRST 10 ROWS ONLY;
```

tagid	tagpath	datatype	floatvalue	dataintegrity	t_stamp	t_stamp_dt
6	sine60	1	32.468379974365234	192	1707780633578	2024-02-12 17:30:33-06
6	sine60	1	27.660436630249023	192	1707780632577	2024-02-12 17:30:32-06
6	sine60	1	23.102643966674805	192	1707780631576	2024-02-12 17:30:31-06
6	sine60	1	18.840131759643555	192	1707780630574	2024-02-12 17:30:30-06
6	sine60	1	14.919697761535645	192	1707780629573	2024-02-12 17:30:29-06
6	sine60	1	11.384378433227539	192	1707780628572	2024-02-12 17:30:28-06
6	sine60	1	8.272985458374023	192	1707780627572	2024-02-12 17:30:27-06
6	sine60	1	5.6196770668029785	192	1707780626571	2024-02-12 17:30:26-06
6	sine60	1	3.4516685009002686	192	1707780625570	2024-02-12 17:30:25-06
6	sine60	1	1.7970842123031616	192	1707780624569	2024-02-12 17:30:24-06

(10 rows)

Tags history statistical (test for Ignition Programmable Device Simulator):

```
tsdb=# select distinct t1.tagid, t2.tagpath, AVG(t1.floatvalue), PERCENTILE_CONT(0.5) WITHIN GROUP (ORDER BY
t1.floatvalue) as median ,MIN(t1.floatvalue), MAX(t1.floatvalue), count(*) FROM public.sqlth_1_data t1,
public.sqlth_te t2 where t1.tagid = t2.id AND t1.dataintegrity =192 group by t1.tagid, t2.tagpath order by t1.tagid;
```

tagid	tagpath	avg	median	min	max	count
1	sine10	49.99550888854456	50	0	100	98046
2	sine20	49.99418173383984	49.952877044677734	0	100	97742
3	sine30	50.002804217046375	50.22514533996582	0	100	97166
4	sine40	50.011398382560586	50.46337890625	0	100	95935
5	sine50	50.00126923737531	49.27745819091797	3.947841946683184e-07	100	95053
6	sine60	50.00151066182858	50.13351821899414	0	100	94072
7	ramp10	52.32011318169635	56.5	0	99.98999786376953	37461
8	ramp20	51.91464345282335	55.27250099182129	0	99.99500274658203	17670
9	ramp30	52.45803161829689	57.74333190917969	0	99.99666595458984	12147
10	ramp40	51.09869536130902	54.3387508392334	0	99.99749755859375	9378
11	ramp50	51.26109527409996	53.57699966430664	0	99.99800109863281	7620
12	ramp60	50.43747796853882	49.310001373291016	0	99.99666595458984	6391

(12 rows)

View Data of Historian DB Infrastructure Timescale Tables and Re-Initialize
To Re-Initialize execute each commented the "delete" command individually:
Note! Ignition historical provider that uses this database must be disabled.

```
-- Delete Tags Historical Records
select count(*) from public.sqlth_1_data;
--delete from public.sqlth_1_data;
--
-- Active/Retired Ignition Tags Historian Configuration
select id, tagpath, scid, datatype, querymode, TO_TIMESTAMP(created/1000), TO_TIMESTAMP(retired/1000) from
public.sqlth_te;
--delete from public.sqlth_te;
--
--Ignition Tags & their Historian Group(s)
select * from public.sqlth_scinfor;
--delete from public.sqlth_scinfor;
--
--Historian driver(s) for Participated Hosts(name) & Tag Provider(s)
select * from public.sqlth_drv;
--delete from public.sqlth_drv;
--
--Historian driver(s) timeframe & rate
select scid, TO_TIMESTAMP(start_time/1000), TO_TIMESTAMP(end_time/1000), rate from public.sqlth_sce;
--delete from public.sqlth_sce;
--
```



```

-- Historian data partitions (tables) & drivers references
select pname, drvid, TO_TIMESTAMP(start_time/1000), TO_TIMESTAMP(end_time/1000), blocksize, flags from
public.sqlth_partitions;
--delete from public.sqlth_partitions;
--
-- annotations
select * from public.sqlth_annotations;
select * from public.saved_graphs;
select * from public.chart_annotations;
--delete from public.sqlth_annotations;
--delete from public.saved_graphs;
--delete from public.chart_annotations;

```

In case if chart_annotations and/or saved_graphs table(s) did not get created (this behaviour was observed):

```

-- DROP SEQUENCE IF EXISTS public.chart_annotations_seq;
CREATE SEQUENCE IF NOT EXISTS public.chart_annotations_seq
    INCREMENT 1
    START 1
    MINVALUE 1
    MAXVALUE 9223372036854775807
    CACHE 1;

ALTER SEQUENCE public.chart_annotations_seq
    OWNER TO ignition;

-- DROP TABLE IF EXISTS public.chart_annotations;
CREATE TABLE IF NOT EXISTS public.chart_annotations
(
    id integer NOT NULL DEFAULT nextval('chart_annotations_seq'::regclass),
    penname character varying(45) COLLATE pg_catalog."default" NOT NULL,
    xvalue timestamp(0) without time zone NOT NULL,
    yvalue double precision NOT NULL,
    note text COLLATE pg_catalog."default" NOT NULL,
    enteredby character varying(45) COLLATE pg_catalog."default" NOT NULL,
    lastupdated timestamp(0) without time zone NOT NULL,
    CONSTRAINT chart_annotations_pkey PRIMARY KEY (id)
) TABLESPACE pg_default;

ALTER TABLE IF EXISTS public.chart_annotations
    OWNER to ignition;

-- DROP INDEX IF EXISTS public.idx_tstamp;
CREATE INDEX IF NOT EXISTS idx_tstamp
    ON public.chart_annotations USING btree
    (xvalue ASC NULLS LAST) TABLESPACE pg_default;
-- *****
-- DROP SEQUENCE IF EXISTS public.saved_graphs_seq;
CREATE SEQUENCE IF NOT EXISTS public.saved_graphs_seq
    INCREMENT 1
    START 1
    MINVALUE 1
    MAXVALUE 9223372036854775807
    CACHE 1;

ALTER SEQUENCE public.saved_graphs_seq
    OWNER TO ignition;

-- DROP TABLE IF EXISTS public.saved_graphs;
CREATE TABLE IF NOT EXISTS public.saved_graphs
(
    id integer NOT NULL DEFAULT nextval('saved_graphs_seq'::regclass),
    title character varying(45) COLLATE pg_catalog."default" NOT NULL,
    tagpens text COLLATE pg_catalog."default" NOT NULL,
    axes text COLLATE pg_catalog."default" NOT NULL,
    chartmode integer,
    startdate timestamp(0) without time zone,
    enddate timestamp(0) without time zone,
    rangestartdate timestamp(0) without time zone,
    rangeenddate timestamp(0) without time zone,
    unit integer,
    unitcount integer,
    username character varying(45) COLLATE pg_catalog."default",
    lastmodified timestamp(0) without time zone NOT NULL,
    CONSTRAINT saved_graphs_pkey PRIMARY KEY (id)
) TABLESPACE pg_default;

ALTER TABLE IF EXISTS public.saved_graphs
    OWNER to ignition;

```

APPENDIX B

Useful queries for PostgreSQL Timescale features and maintenance

View Autovacuum activities:

```
tsdb=# SELECT relname, last_vacuum, last_autovacuum FROM pg_stat_user_tables
WHERE last_autovacuum is NOT NULL ORDER BY last_autovacuum DESC;
 relname      | last_vacuum | last_autovacuum
-----|-----|-----
sqlth_sce     |             | 2024-02-12 17:45:17.738351-06
_hyper_1_5_chunk |           | 2024-02-12 16:47:24.420229-06
_hyper_1_4_chunk |           | 2024-02-11 01:09:40.624457-06
_hyper_1_3_chunk |           | 2024-02-10 15:38:59.95357-06
bgw_job_stat  |           | 2024-02-09 20:07:52.350248-06
_hyper_1_2_chunk |           | 2024-02-09 17:43:39.675193-06
(6 rows)
```

View specific database size:

```
tsdb=# SELECT pg_size_pretty(pg_database_size('tsdb'));
pg_size_pretty
-----
73 MB
(1 row)
```

Get the disk space used by an index on a hypertable, including the disk space needed to provide the index on all chunks. The size is reported in bytes:

```
tsdb=# SELECT hypertable_index_size('sqlth_1_data_tagid_t_stamp_brinx');
hypertable_index_size
-----
147456
(1 row)
```

Show hypertable chunks:

```
tsdb=# SELECT show_chunks('sqlth_1_data');
show_chunks
-----
_timescaledb_internal._hyper_1_1_chunk
_timescaledb_internal._hyper_1_2_chunk
_timescaledb_internal._hyper_1_3_chunk
_timescaledb_internal._hyper_1_4_chunk
_timescaledb_internal._hyper_1_5_chunk
(5 rows)
```

Maintenance jobs setup and modifications

```
-- ==== Maintenance JOBS ====
-- Show all jobs for timescaledb:
SELECT * FROM timescaledb_information.jobs;
--
-- ==== HYPERTABLE RETENTION POLICY ====
-- to remove JOB for retention policy on a given 'sqlth_X_data' hypertable
SELECT remove_retention_policy('sqlth_1_data');
--
-- to add retention policy to 'sqlth_X_data' hypertable to drop chunks older than:
-- 30 days (2592000000 ms), OR
-- 366 days (31622400000 ms)
SELECT add_retention_policy('sqlth_1_data', BIGINT '2592000000');
--
-- ==== HYPERTABLE COMPRESSION POLICY =====
-- to remove JOB for compression policy on a given 'sqlth_X_data' hypertable
SELECT remove_compression_policy('sqlth_1_data');
--
-- to add compression policy to 'sqlth_X_data' hypertable for chunks age more than:
-- 7 days (604800000 ms), OR
-- 30 days (2592000000 ms), OR
-- 31 days (2678400000 ms)
SELECT add_compression_policy('sqlth_1_data', BIGINT '604800000');
```


Ignition Tag Historian Module:

<https://inductiveautomation.com/resources/article/scada-historian>

BRIN Indexes by PostgreSQL:

<https://www.postgresql.org/docs/current/brin-intro.html>

BRIN Index for PostgreSQL by PERCONA:

<https://www.percona.com/blog/brin-index-for-postgresql-dont-forget-the-benefits/>

Optimizing Postgres's Autovacuum for High-Churn Tables by Adam Hendel:

<https://tembo.io/blog/optimizing-postgres-auto-vacuum>

Enabling and Disabling Autovacuum in Postgresql by Hans-Jürgen Schöning:

<https://www.cybertec-postgresql.com/en/enabling-and-disabling-autovacuum-in-postgresql/>

Postgres Indexing: When Does BRIN Win? By Paul Ramsey:

<https://www.crunchydata.com/blog/postgres-indexing-when-does-brin-win>

Create an BRIN Index at a Fraction of the Normal Size by Digoal Zhou:

https://www.alibabacloud.com/blog/create-an-brin-index-at-a-fraction-of-the-normal-size_595138

Install self-hosted TimescaleDB on Linux:

<https://docs.timescale.com/self-hosted/latest/install/installation-linux/>

{EOF}